

A GPU Accelerated Algorithm for 3D Delaunay Triangulation



**THANH-TUNG CAO, TODD
MINGCEN GAO
TIOW-SENG TAN**

*National University of
Singapore*

ASHWIN NANJAPPA

*Bioinformatics Institute
Singapore*


Outline

→ Background


2

- Background
- Related work
- Algorithm
- Implementation
- Result

- Delaunay Triangulation
 - Empty ball property



2D




3D

- Applications:
 - Graphics
 - CAD
 - Visualization
 - Scientific computation



- Sequential algorithms:
 - Incremental construction
 - Divide-and-conquer
 - Incremental insertion
 - ✦ Points are inserted one by one
 - ✦ Triangulation is locally fixed after each insertion
- Bowyer-Watson's algorithm [1981]
- Flipping algorithm [Joe 1991]



Related work


6

- **Parallel and multi-core algorithms:**
 - Incremental construction → high complexity
 - Domain partitioning → GPU needs thousands of partitions
- Incremental insertion [Batista et al. 2010]
 - ✦ Several points are inserted at a time
 - ✦ Each insertion modifies a small region
 - ✦ Conflict → Rollback

- GPU algorithms:
 - Experiment with Batista et al.'s approach
 - ✦ 1 million points
 - ✦ At most 2000-3000 points can be inserted in each round
 - Digital Voronoi diagram approach [Qi et al. 2012]
 - ✦ Dualization not working in 3D

- Overall approach:

- Insert points in batches, each batch at most one point is inserted into a tetrahedron
- Use flipping to get close to the DT
- Use star-splaying in CPU to fix [Shewchuk 2005]
- Problem: Flipping easily gets stuck!
 - ✦ 100K points, $\approx 6,800$ bad facets (non-Delaunay and unflippable).
 - ✦ Worse for real-world data.



Algorithm

9

- Observation:

- Do flipping after each round of point insertion

- Much better result.

- ✦ 100K points, 96 bad facets (vs. 6,800)

- ✦ Bunny model (36K points), 92 bad facets


- ✦ Now correction on CPU is acceptable.



Algorithm

10

- CPU correction: star splaying algorithm
 - Lifting map: $w = x^2 + y^2 + z^2$
 - Each vertex: construct a convex star.
 - Consistency: If the star of s contains tetrahedron $stuv$, then the star of t , u , and v must also contain that tetrahedron.




2D illustration

Algorithm

11


- CPU correction: star splaying algorithm
 - Difficulties: Time consuming
 - ✦ Construct convex stars (incremental insertion)
 - ✦ Check all the stars for inconsistencies
 - ✦ Convert stars back to mesh representation



Algorithm

12

- Adaptive star splaying
 - Only some small regions around the bad facets are processed
 - ✦ Construct stars incident to the bad facets.
 - ✦ Almost convex → use Flip-Flop [Gao et al. 2013]
 - ✦ Need another star → derive from the triangulation.
 - Almost output sensitive.



2D illustration

Initialize;

While there are points not inserted

Pick one point per tetrahedron and insert;

While there are modified tetrahedra

Collect the modified tetrahedra;

Process and identify possible flips;

Perform the flips;

Update location of remaining points;

GPU Implementation

14

- Thread divergence:
 - Compact the list of modified tetrahedra before processing
 - Exact predicates [Shewchuk 1996]
 - ✦ Use only fast check in 1st kernel
 - ✦ Collect those that require exact computation
 - ✦ Do the exact computation in 2nd kernel
 - Point location:
 - ✦ Store the flips, build flip history DAG
 - ✦ Trace the DAG to locate

- **Memory access:**
 - Rearrange the data to improve the GPU cache efficiency.
 - ✦ Sort the input points by the Z-curve
 - ✦ Sort the tetrahedra list by the minimum vertex indices.

- Experiment settings:
 - CPU: Intel I7 2600K 3.4Ghz, 16GB RAM
 - GPU: NVIDIA GTX 580, 3GB VRAM
 - CUDA 5.0, VS 2012.
 - CGAL 4.2

Result

17

- 3D Speedup:
 - Synthetic data: Uniform, Gaussian, sphere, grid.
 - ✦ Up to 1.5 million points
 - ✦ 8-10 times faster than CGAL
 - Real models: Armadillo, Angel, Dragon, Happy Buddha...
 - ✦ 6-9 times speedup over CGAL

Result


18

- Also implement for 2D DT:
 - Synthetic data:
 - ✦ 10 times over *Triangle*, 7 times over CGAL
 - ✦ 2 times faster than [Qi et al. 2012]

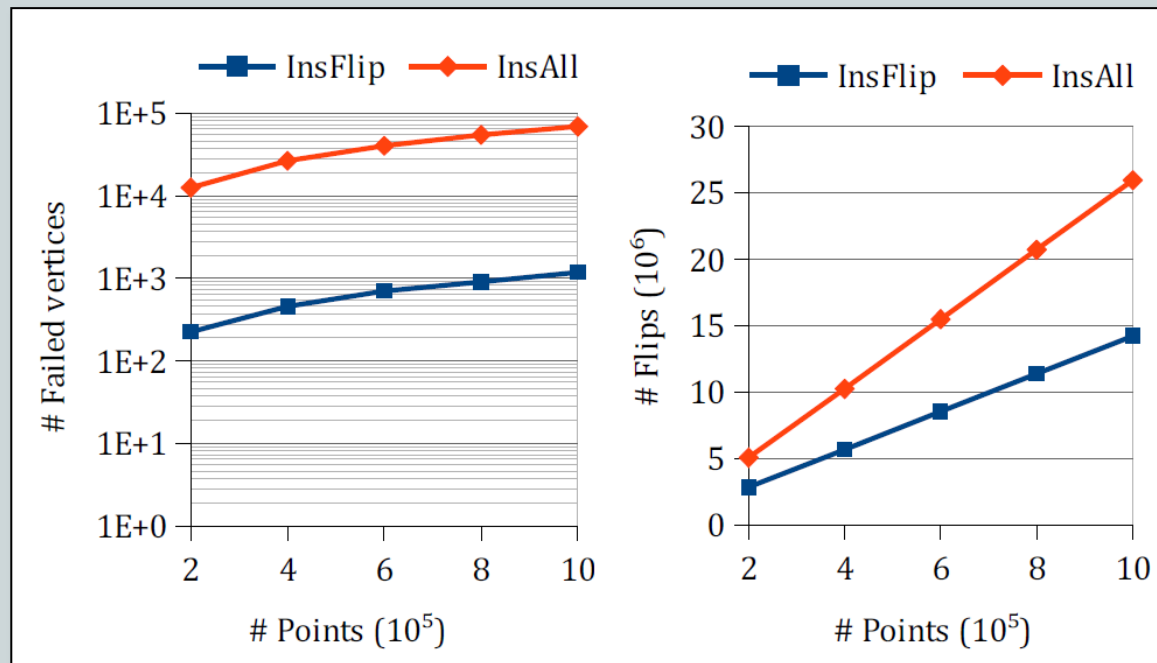
Result

19

- Time breakdown



- Insert-Flip vs. InsertAll-Flip
 - 100 times less bad facets
 - 40% less flips



- New algorithm for DT construction on GPU
- Both 2D and 3D (possibly higher)
- Uniform and non-uniform point set
- Exact computation, robust.
- Limitation:
 - Needs to copy the result to CPU for splaying
 - Sequential flipping on some pathological cases
 - Memory bound implementation

End

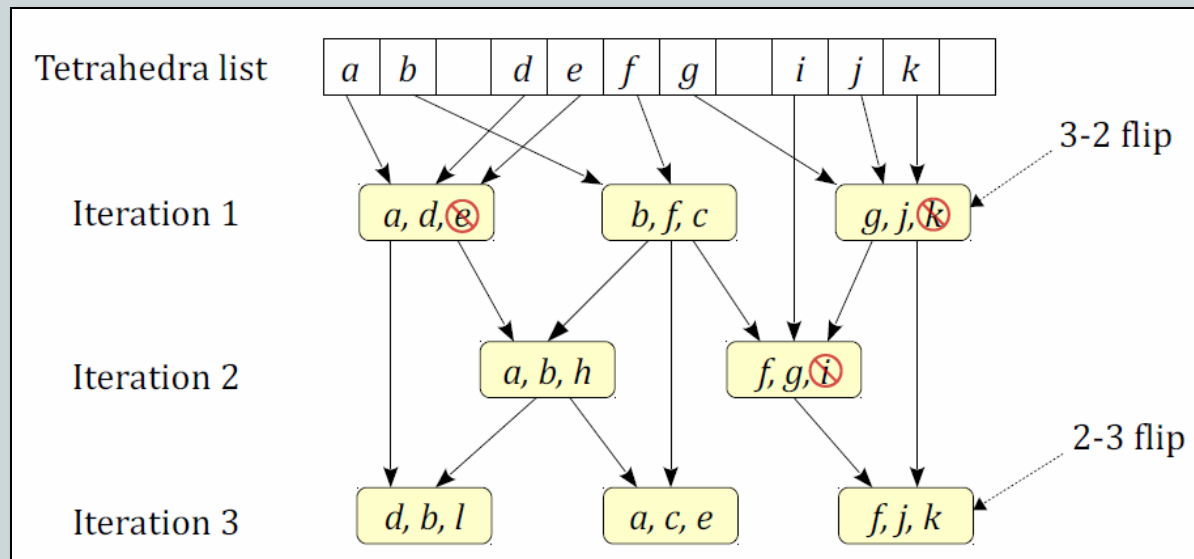
22

Thank you!

GPU Implementation

23

- Point location:
 - Store the flips in all the iterations
 - Construct the history DAG of flips
 - Update the point location at the end



Result

24

- Stars involved in the CPU star splaying
 - Significantly more for non-uniform point sets
 - Still reasonably small

